

# Northeastern University – Silicon Valley

## CS 6620 Cloud Computing

### Homework Set #5 [100 points]

---

***INSTRUCTIONS:*** Please provide clear explanations in your own sentences, directly answering the question, demonstrating your understanding of the question and its solution, in depth, with sufficient detail. Submit your solutions [PDF preferred]. Include your full name. Do not email the solutions.

#### **PART I: Concepts and Theory, Algorithms [60 points]**

Please provide accurate, concise (maximum 10 sentence) answers. All questions carry 10 points.

1. Based on your study of text/references, explain (i) what each of the below AWS feature is and (ii) why and how they are needed/used in one example business application (Chapter 5): [20 points]
  - A. Amazon S3
  - B. Amazon EFS
  - C. Amazon EBS
  - D. AWS Storage Gateway

Why are these 4 separate services necessary? That is, how are they different?

2. Explain what a Storage Area network (SAN) and Network Attached Storage (NAS) are, and their differences [10 points]
3. Explain the purpose of and differences among a database, data lake, data warehouse. Which AWS storage services are suitable to build each of these on the cloud? [10 points]
4. Briefly explain how SSD, HDD, Flash disk and Optical Disk are different. Separately, explain what storage Tiering is and why it is needed in Cloud. Then, identify which storage technology (SSD etc) is suitable for which tier. (See references). [15 points]

#### **PART II: LAB [45 points]**

**Coding needed:** Following the examples in Chapter 5 of the textbook OR external links and examples, implement a basic setup and use of the below 3 AWS storage services. Goal is to demonstrate the basic functionality only:

- A. Amazon S3
- B. Amazon EFS
- C. Amazon EBS

Advanced features like security, ACLS etc are not necessary for this exercise.

#### **References**

<https://cloudian.com/guides/data-backup/storage-tiering/>  
<https://experience.dropbox.com/get-organized/storage-devices>  
<https://www.sqlshack.com/getting-started-with-amazon-s3-and-python/>

## PART I: Concepts and Theory, Algorithms

1. Based on your study of text/references, explain (i) what each of the below AWS feature is and (ii) why and how they are needed/used in one example business application (Chapter 5): [20 points]

### A. Amazon S3

For the concept of Amazon S3 (Amazon Simple Storage Service), according to the website of amazon, it is an object storage service that offers industry-leading scalability, data availability, security, and performance. This means customers of all sizes and industries can use it to store and protect any amount of data for a range of use cases, such as data lakes, websites, cloud-native applications, backups, archive, machine learning, and analytics.

For the explanation of examples, Amazon S3 could build fast, powerful mobile and web-based cloud-native apps that scale automatically in a highly available configuration while running Cloud-Native Applications.

### B. Amazon EFS

For the concept of Amazon Elastic File System (Amazon EFS), according to the website of amazon, it provides a simple, scalable, elastic file system for Linux-based workloads for use with AWS Cloud services and on-premises resources. It is built to scale on demand to petabytes without disrupting applications, growing and shrinking automatically as you add and remove files, so your applications have the storage they need – when they need it. It is designed to provide massively parallel shared access to thousands of Amazon EC2 instances, enabling your applications to achieve high levels of aggregate throughput and IOPS with consistent low latencies. Amazon EFS is a fully managed service that requires no changes to your existing applications and tools, providing access through a standard file system interface for seamless integration. Amazon EFS is a regional service storing data within and across multiple Availability Zones (AZs) for high availability and durability. You can access your file systems across AZs and AWS Regions and share files between thousands of Amazon EC2 instances and on-premises servers via AWS Direct Connect or AWS VPN.

For the explanation of examples -- machine learning (ML) and big data analytics workloads, Amazon EFS could accelerate data science and then it's easy to use and scale, Amazon EFS offers the performance and consistency.

### C. Amazon EBS

For the concept of Amazon Elastic Block Store (Amazon EBS), it provides persistent block storage volumes for use with Amazon EC2 instances in the AWS Cloud. Each Amazon EBS volume is automatically replicated within its Availability Zone to protect you from component failure, offering high availability and durability. Amazon EBS volumes offer the consistent and low-latency performance needed to run your workloads. With Amazon EBS, you can scale your usage up or down within minutes—all while paying a low price for only what you provision.

For the explanation of examples, we should know that Amazon EBS provides the following volume types: General Purpose SSD, Provisioned IOPS SSD, Throughput Optimized HDD, and Cold HDD. For the general Purpose SSD volumes (gp2 and gp3), it balances price and performance for a wide variety of transactional workloads. Because these volumes are ideal for use cases such as boot volumes, medium-size single instance databases, and development and test environments.

### D. AWS Storage Gateway

For the concept of the Storage Gateway, according to the website of amazon, it is a hybrid storage service that enables your on-premises applications to seamlessly use AWS cloud storage. You can use the service for backup and archiving, disaster recovery, cloud data processing, storage tiering, and migration. Your applications connect

to the service through a virtual machine or hardware gateway appliance using standard storage protocols, such as NFS, SMB and iSCSI. The gateway connects to AWS storage services, such as Amazon S3, S3 Glacier, and Amazon EBS, providing storage for files, volumes, and virtual tapes in AWS. The service includes a highly optimized data transfer mechanism, with bandwidth management, automated network resilience, and efficient data transfer, along with a local cache for low-latency on-premises access to your most active data.

For the explanation of examples, the Storage Gateway supports four key hybrid cloud use cases – (1) move backups and archives to the cloud, (2) reduce on-premises storage with cloud-backed file shares, (3) provide on-premises applications low latency access to data stored in AWS, and (4) data lake access for pre and post processing workflows. And depending on the use case, Storage Gateway provides three types of storage interfaces for the on-premises applications: file, volume, and tape.

Why are these 4 separate services necessary? That is, how are they different?

The four services work quite differently and offer different levels of performance, cost, availability, and scalability. The table below from online website compares Amazon S3, EBS, and EFS in terms of performance, cost, availability, accessibility, access control, and storage or file size limits enforced by each service.

	Performance	Availability and Accessibility	Access Control	Storage and File Size Limits	Cost
<b>Amazon S3</b>	<ul style="list-style-type: none"> <li>- Supports 3500 PUT / LIST / DELETE requests per second</li> <li>- Scalable to 5500 GET requests per second</li> </ul>	<ul style="list-style-type: none"> <li>- Usually 99.9% available</li> <li>- If lower, returns 10-100% of cost as service credits</li> <li>- Accessible via Internet using APIs</li> </ul>	<ul style="list-style-type: none"> <li>- Access is based on IAM</li> <li>- Uses bucket policies and user policies</li> <li>- Public access via Block Public Access</li> </ul>	<ul style="list-style-type: none"> <li>- No limit on quantity of objects</li> <li>- Individual objects up to 5TB</li> </ul>	<ul style="list-style-type: none"> <li>- Free tier: 5GB</li> <li>- First 50 TB/month: \$0.023 per GB</li> <li>- Next 450 TB/month: \$0.022 per GB</li> <li>- Over 500 TB/month: \$0.021 per GB</li> </ul>
<b>AWS EBS</b>	<ul style="list-style-type: none"> <li>- HDD volumes: 250-500 IOPS/volume depending on volume type</li> <li>- SSD volumes: 16-64K IOPS/volume</li> </ul>	<ul style="list-style-type: none"> <li>- 99.99% available</li> <li>- Accessible via single EC2 instance</li> </ul>	<ul style="list-style-type: none"> <li>- Security groups</li> <li>- User-based authentication (IAM)</li> </ul>	<ul style="list-style-type: none"> <li>- Max storage size of 16TB</li> <li>- No file size limit on disk</li> </ul>	<ul style="list-style-type: none"> <li>- Free tier: 30GB</li> <li>- General Purpose: \$0.045 per GB/month</li> <li>- Provisioned SSD: \$0.125 per GB/month, \$0.065 per IOPS/month</li> </ul>
<b>AWS EFS</b>	<ul style="list-style-type: none"> <li>- 3GB/s baseline performance</li> <li>- Up to 10GB/s</li> <li>- Up to 7K IOPS</li> </ul>	<ul style="list-style-type: none"> <li>- No publicly available SLA</li> <li>- Up to 1,000 concurrent EC2 instances</li> <li>- Accessible from any AZ or region</li> </ul>	<ul style="list-style-type: none"> <li>- IAM user-based authentication</li> <li>- Security groups</li> </ul>	<ul style="list-style-type: none"> <li>- 16TB per volume</li> <li>- 52TB maximum for individual files</li> </ul>	<ul style="list-style-type: none"> <li>- Standard storage: \$0.30-\$0.39 per GB-month depending on region</li> <li>- Infrequent storage: \$0.025-\$0.03 per GB-month</li> <li>- Provisioned throughput: \$6 per MB/s-month</li> </ul>

In detail:

Amazon S3 is cheapest for data storage alone. However, there are various other pricing parameters in S3, including cost per number of requests made, S3 Analytics, and data transfer out of S3 per gigabyte. EFS has the simplest cost structure.

Amazon S3 can be accessed from anywhere. AWS EBS is only available in a particular region, while you can share files between regions on multiple EFS instances.

EBS and EFS are both faster than Amazon S3, with high IOPS and lower latency.

EBS is scalable up or down with a single API call. Since EBS is cheaper than EFS, you can use it for database backups and other low-latency interactive applications that require consistent, predictable performance.

EFS is best used for large quantities of data, such as large analytic workloads. Data at this scale cannot be stored on a single EC2 instance allowed in EBS—requiring users to break up data and distribute it between EBS instances. The EFS service allows concurrent access to thousands of EC2 instances, making it possible to process and analyze large amounts of data seamlessly.

Storage Gateway is a Hybrid Cloud for Storage. Part of the infrastructure is on the cloud; another part is on-premises. EBS is block, EFS is file, S3 is object. Whereas Storage Gateway are fancy local caches and gateways to access the remote storage with familiar protocols.

Storage Gateway still requires the other storage systems. If AWS is like a non-cloud remote storage array, SG is the filer and VTL options that speak file and tape.

## 2. Explain what a Storage Area network (SAN) and Network Attached Storage (NAS) are, and their differences [10 points]

For the concept of Storage Area Network (SAN), it is a specialized, high-speed network that provides block-level network access to storage. SANs are typically composed of hosts, switches, storage elements, and storage devices that are interconnected using a variety of technologies, topologies, and protocols. SANs may also span multiple sites.

A SAN presents storage devices to a host such that the storage appears to be locally attached. This simplified presentation of storage to a host is accomplished through the use of different types of virtualization.

For the concept of Network-attached storage (NAS), it is a file-level (as opposed to block-level storage) computer data storage server connected to a computer network providing data access to a heterogeneous group of clients. NAS is specialized for serving files either by its hardware, software, or configuration.

Differences:

NAS is a single storage device that serves files over Ethernet and is relatively inexpensive and easy to set up, while a SAN is a tightly coupled network of multiple devices that is more expensive and complex to set up and manage. From a user perspective, the biggest difference between NAS and SAN is that NAS devices deliver shared storage as network mounted volumes and use protocols like NFS and SMB/CIFS, while SAN-connected disks appear to the user as local drives.

## 3. Explain the purpose of and differences among a database, data lake, data warehouse. Which AWS storage services are suitable to build each of these on the cloud? [10 points]

Typically, an organization will require a data lake, data warehouse and database(s) for different use cases.

A database is used to store, search and report on structured data from a single source. They are the simplest to create and SQL can be used to query and report on the data.

A data warehouse is used to store large amounts of structured data from multiple sources in a centralized place. Organizations invest in building data warehouses because of its ability to deliver business insights from across the company, and quickly.

A data lake stores structured, semi-structured and unstructured data, supporting the ability to store raw data from all sources without the need to process or transform it at that time.

For the difference, a data lake is where streams and rivers of data from various sources meet. All data is allowed, no matter if it is structured or unstructured and no processing is done to the data until after it is in the data lake. However, a data warehouse is a centralized place for structured data to be analyzed for specific purposes related to business insights. The requirements for reporting are known ahead of time during the planning and design of a data warehouse and the ETL process. For a database, it thrives in a monolithic environment where the data is being generated by one application.

Related AWS storage services:

Amazon S3/ Amazon Redshift -- data lake;

Amazon Redshift -- data warehouse;

Amazon RDS/ Amazon DynamoDB -- database

4. Briefly explain how SSD, HDD, Flash disk and Optical Disk are different. Separately, explain what storage Tiering is and why it is needed in Cloud. Then, identify which storage technology (SSD etc) is suitable for which tier. (See references). [15 points]

Both HDD and SSD devices generally offer the largest storage capacity among external options, with external HDDs offering up to 20 TB of storage and external SSDs offering up to 8 TB of storage.

For flash memory devices, one of the most recognizable type of flash memory device is the USB flash drive, which can hold up to 2 TB of storage.

For Optical Disk, CD can store up to 700 MB of data, DVD-DL can store up to 8.5 GB, and Blu-Ray can store between 25 and 128 GB of data.

Storage tiering is a strategy that lets you optimize the use of storage resources, efficiently backup data, save costs and make the best use of storage technology for each data class.

Why it is needed: for example, SSD disk drives, magnetic disk drives and tape storage. The most important or frequently-accessed data is stored on the fastest, and most expensive media (SSD) and the least important on the slowest, cheapest media (tape). The minimal storage tiering system has two tiers—one for frequently accessed data and one for archive. The more tiers are available, the more choices administrators have over the placement of specific data classes, and the more efficiently storage resources can be utilized.

As result, in Cloud, storage Tiering can reduce public cloud storage costs.

			Type	Storage Media	Used For
			Tier 0	Tier 0 includes SSD, RAM, PCIe Flash	You can use Tier 0 for high performance workloads.
			Tier 1	Tier 1 includes fast disks, all-flash storage, hybrid flash storage	You can use Tier 1 for mission-critical or highly sensitive files.
	<b>Cold</b>	<b>Hot</b>			
<b>Required Access Speed</b>	Slow	Fast			
<b>Access Frequency</b>	Low	High	Tier 2/3	Tier 2 and Tier 3 include Slow-spinning HDD, disk-based backup appliance, cloud storage, tape	You can use Tier 2 and Tier 3 for backups of mission critical data, which requires high reliability but not instant retrieval from backup.
<b>Value of Data</b>	Low	High			
<b>Storage Media</b>	Slower drives, tape	Faster drives, SSD	Tier 4	Tier 4 include SATA drives	You can use Tier 4 for warm data, data used for periodic reporting.
<b>Storage Location</b>	May be off-premises	Colocated or fast link to the data consumer			
<b>Cost</b>	Low cost	High cost	Tier 5	Tier 5 includes tape storage, cloud storage archive tiers (e.g. Amazon Glacier)	You can use Tier 5 for cold data which is rarely or never accessed.

SSD is suitable for fast and high-performance tier and tape is suitable for slow and low cost tier.

## PART II: LAB [45 points]

### A. Amazon S3

Configuration:  
Assigned role with S3 read/write permission

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

Add user

12345

Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name\*

test-recognition

Add another user

Select AWS access type

Select how these users will primarily access AWS. If you choose only programmatic access, it does NOT prevent users from accessing the console using an assumed role. Access keys and autogenerated passwords are provided in the last step. [Learn more](#)

Select AWS credential type\*

☒ Access key - Programmatic access

Enables an **access key ID** and **secret access key** for the AWS API, CLI, SDK, and other development tools.

☐ Password - AWS Management Console access

Enables a **password** that allows users to sign-in to the AWS Management Console.

\* Required

Cancel

Next: Permissions

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

Add user

12345

Set permissions

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

rekog

Showing 4 results

	Policy name	Type	Used as
<input type="checkbox"/>	AmazonRekognitionCustomLabelsFullAccess	AWS managed	None
<input checked="" type="checkbox"/>	AmazonRekognitionFullAccess	AWS managed	None
<input type="checkbox"/>	AmazonRekognitionReadOnlyAccess	AWS managed	None
<input type="checkbox"/>	AmazonRekognitionServiceRole	AWS managed	None

Set permissions boundary

Cancel

Previous

Next: Tags

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

7

aws

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

Add user

12345

Set permissions

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

s3

Showing 7 results

	Policy name	Type	Used as
<input type="checkbox"/>	AmazonDMSRedshiftS3Role	AWS managed	None
<input checked="" type="checkbox"/>	AmazonS3FullAccess	AWS managed	Permissions policy (1)
<input type="checkbox"/>	AmazonS3ObjectLambdaExecutionRolePolicy	AWS managed	None
<input type="checkbox"/>	AmazonS3OutpostsFullAccess	AWS managed	None
<input type="checkbox"/>	AmazonS3OutpostsReadOnlyAccess	AWS managed	None
<input type="checkbox"/>	AmazonS3ReadOnlyAccess	AWS managed	None
<input type="checkbox"/>	QuickSightAccessForS3StorageManagementAnalyticsReadOnly	AWS managed	None

Set permissions boundary

CancelPreviousNext: Tags

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

aws

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

Add user

12345

Review

Review your choices. After you create the user, you can view and download the autogenerated password and access key.

User details

User name

test-recognition

AWS access type

Programmatic access - with an access key

Permissions boundary

Permissions boundary is not set

Permissions summary

The following policies will be attached to the user shown above.

Type	Name
Managed policy	AmazonRekognitionFullAccess
Managed policy	AmazonS3FullAccess

Tags

No tags were added.

CancelPreviousCreate user

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

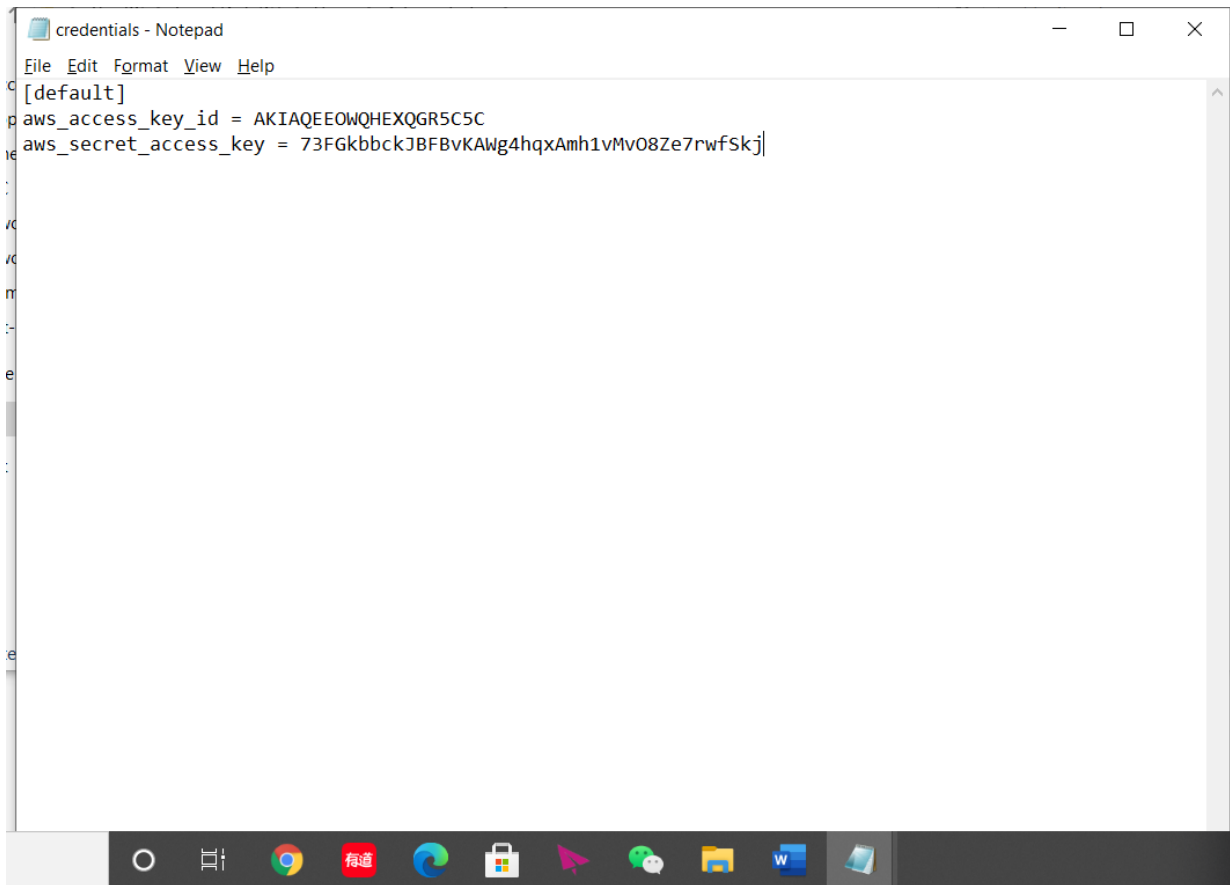
Privacy Policy

Terms of Use

Cookie preferences



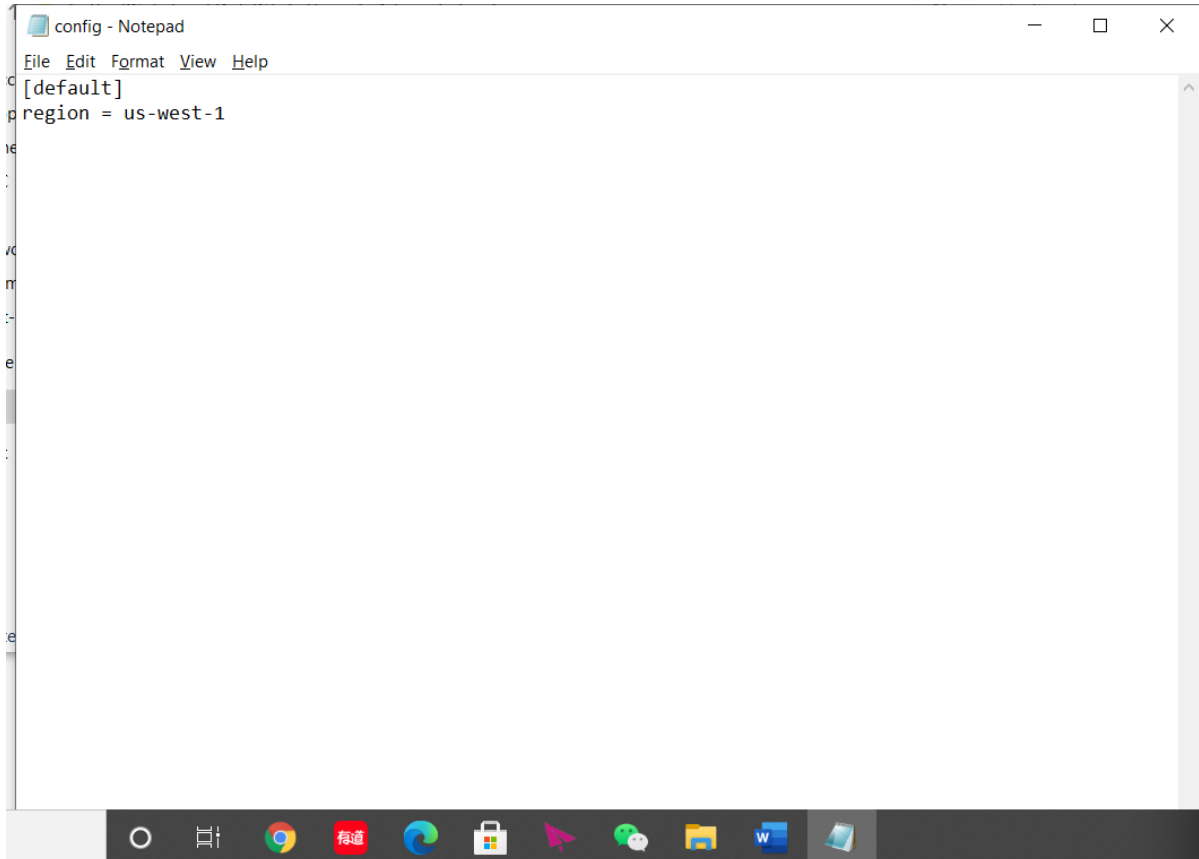
Access\_key:



A screenshot of a Windows Notepad application window titled "credentials - Notepad". The window has a standard menu bar with "File", "Edit", "Format", "View", and "Help". The text content is as follows:

```
[default]
aws_access_key_id = AKIAQEEOWQHEXQGR5C5C
aws_secret_access_key = 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]
```

The window is positioned over a Windows taskbar which includes icons for the Start button, Task View, Google Chrome, WeChat, Edge, File Explorer, and several other applications.



A screenshot of a Windows Notepad application window titled "config - Notepad". The window has a standard menu bar with "File", "Edit", "Format", "View", and "Help". The text content is as follows:

```
[default]
region = us-west-1
```

The window is positioned over a Windows taskbar which includes icons for the Start button, Task View, Google Chrome, WeChat, Edge, File Explorer, and several other applications.

Upload the image that contains text to S3 bucket.

Upload:

```
[root@ip-172-31-10-156 s3]# aws s3 cp ./ s3://project2-image/ --recursive
upload: ./imageFile.jpg to s3://project2-image/imageFile.jpg
```

Library:

```
aqc@ubuntu:~/aqc/lambda/s3$ pip install boto3
aqc@ubuntu:~/aqc/lambda/s3$ pip install pandas
collecting pandas
```

Print a list of all the buckets:

```
aqc@ubuntu: ~/aqc/lambda/s3

# Creating the low level functional client
client = boto3.client(
    's3',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Creating the high level object oriented interface
resource = boto3.resource(
    's3',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Fetch the list of existing buckets
clientResponse = client.list_buckets()

# Print the bucket names one by one
print('Printing bucket names...')
for bucket in clientResponse['Buckets']:

root@ubuntu:/home/aqc/aqc/lambda/s3# python3 s3.py
Printing bucket names...
Bucket Name: codepipeline-us-west-1-971689687851
Bucket Name: project2-image
Bucket Name: zappa-pgk1rhhn5
Bucket Name: zappa-rw3sg9dt9
Bucket Name: zappa-tn4u6j73t
```

## Create a bucket

```
root@ubuntu: /home/aqc/aqc/lambda/s3

# Creating the low level functional client
client = boto3.client(
    's3',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Creating the high level object oriented interface
resource = boto3.resource(
    's3',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Creating a bucket in AWS S3
location = {'LocationConstraint': 'us-west-1'}
client.create_bucket(
    Bucket='hw5-demo-1',
    CreateBucketConfiguration=location
)

-- INSERT --
```

```
root@ubuntu:/home/aqc/aqc/lambda/s3# python3 s3create.py
root@ubuntu:/home/aqc/aqc/lambda/s3# python3 s3.py
Printing bucket names...
Bucket Name: codepipeline-us-west-1-971689687851
Bucket Name: hw5-demo-1
Bucket Name: project2-image
Bucket Name: zappa-pgk1rhhn5
Bucket Name: zappa-rw3sg9dt9
Bucket Name: zappa-tn4u6j73t
```

```

root@ubuntu: /home/aqc/aqc/lambda/s3
's3',
aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
aws_secret_access_key = '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',
region_name = 'us-west-1'
)

# Creating the high level object oriented interface
resource = boto3.resource(
    's3',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Create the S3 object
obj = client.get_object(
    Bucket = 'project2-image',
    Key = 'static/text'
)

# Read data from the S3 object
data = pandas.read_csv(obj['Body'])

"s3read.py" 31L, 752C 23,19 62%
root@ubuntu:/home/aqc/aqc/lambda/s3# python3 s3read.py
Printing the data frame...
Detected text
0
1 Detected text:Be the Author of
2 Confidence: 65.88%
3 Id: 0
4 Type:LINE
..
138 Confidence: 100.00%
139 Id: 29
140 Parent Id: 8
141 Type:WORD
142 Text detected: 30

[143 rows x 1 columns]

```

## B. Amazon EFS

### Configuration:

Assigned role with EFS read/write permission

1 2

### Add permissions to test-recognition

#### Grant permissions

Use IAM policies to grant permissions. You can assign an existing policy or create a new one.

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

Showing 6 results

	Policy name	Type	Used as
<input type="checkbox"/>	AmazonElasticFileSystemClientFullAccess	AWS managed	None
<input type="checkbox"/>	AmazonElasticFileSystemClientReadOnlyAccess	AWS managed	None
<input type="checkbox"/>	AmazonElasticFileSystemClientReadWriteAccess	AWS managed	None
<input checked="" type="checkbox"/>	AmazonElasticFileSystemFullAccess	AWS managed	None
<input type="checkbox"/>	AmazonElasticFileSystemReadOnlyAccess	AWS managed	None
<input type="checkbox"/>	AmazonElasticFileSystemsUtils	AWS managed	None

Cancel Next: Review

Feedback English (US) © 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences

### Create EFS:

```
root@ubuntu: /home/aqc/aqc/lambda/s3

import boto3
import pandas

# Creating the low level functional client
client = boto3.client(
    'efs',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

client.create_file_system(
    CreationToken='efs-demo1',
    Tags=[
        {
            'Key': 'aws',
            'Value': 'efs-demo1'
        }
    ],
)
```

3,0-1 All

aws Services Search for services, features, marketplace products, and docs [Alt+S] scyqa1 N. California Support

Elastic File System X

File systems Access points

AWS Backup AWS DataSync AWS Transfer Documentation

Amazon EFS File systems

Reduce your storage price to \$0.08/GB-month\* using EFS Lifecycle Management and Intelligent-Tiering. [Learn more](#)

\* pricing in US East (N. Virginia) region, assumes 80% of your storage in EFS IA

[What's new](#) | [Documentation](#) | [AWS Storage Blog](#)

File systems (1) View details Delete Create file system

Filter by property values

Name	File system ID	Encryption	Total size	Size in Standard / One Zone	Size in Standard-IA / One Zone-IA	Provisioned Throughput (MiB/s)	File system state	Creation time	Availability Zone
-	fs-05adc37d783b6392d	Unencrypted	6.00 KiB	6.00 KiB	0 Bytes	-	Available	Wed, 27 Oct 2021 05:23:44 GMT	Regional

Feedback English (US) © 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use Cookie preferences

Delete EFS:

```
root@ubuntu: /home/aqc/aqc/lambda/s3
client = boto3.client(
    'efs',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Creating the high level object oriented interface
resource = boto3.resource(
    'efs',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

client.delete_file_system(
    FileSystemId='fs-05adc37d783b6392d'
)

-- INSERT --

root@ubuntu: /home/aqc/aqc/lambda/s3# python3 efsDelete.py
```

C. Amazon EBS

Configuration:  
Assigned role with EBS read/write permission

aws

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

Add permissions to test-recognition

12

Grant permissions

Use IAM policies to grant permissions. You can assign an existing policy or create a new one.

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

storage

Showing 3 results

	Policy name	Type	Used as
<input checked="" type="checkbox"/>	<div>AWSS3OutpostsFullAccess</div>	AWS managed	None
<input type="checkbox"/>	<div>AWSS3OutpostsReadOnlyAccess</div>	AWS managed	None
<input type="checkbox"/>	<div>QuickSightAccessForS3StorageManagementAnalyticsReadOnly</div>	AWS managed	None

Cancel

Next: Review

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

aws

Services

Search for services, features, marketplace products, and docs

[Alt+S]

scyqa1

Global

Support

use IAM policies to grant permissions. you can assign an existing policy or create a new one.

Add user to group

Copy permissions from existing user

Attach existing policies directly

Create policy

Filter policies

ec2

Showing 25 results

	Policy name	Type	Used as
<input type="checkbox"/>	<div>AmazonEC2ContainerRegistryFullAccess</div>	AWS managed	Permissions policy (2)
<input type="checkbox"/>	<div>AmazonEC2ContainerRegistryPowerUser</div>	AWS managed	Permissions policy (1)
<input type="checkbox"/>	<div>AmazonEC2ContainerRegistryReadOnly</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2ContainerServiceAutoscaleRole</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2ContainerServiceEventsRole</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2ContainerServiceforEC2Role</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2ContainerServiceRole</div>	AWS managed	None
<input checked="" type="checkbox"/>	<div>AmazonEC2FullAccess</div>	AWS managed	Permissions policy (1)
<input type="checkbox"/>	<div>AmazonEC2ReadOnlyAccess</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2RoleforAWSCodeDeploy</div>	AWS managed	None
<input type="checkbox"/>	<div>AmazonEC2RoleforAWSCodeDeployLimited</div>	AWS managed	None

Cancel

Next: Review

Feedback

English (US)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use

Cookie preferences

Create snapshot:

The image displays a terminal window and the AWS Management Console. The terminal shows a Python script for creating an EBS snapshot using boto3. The script imports boto3 and pandas, creates a boto3 client and resource for the 'ebs' service in the 'us-west-1' region, and then calls client.start\_snapshot() with VolumeSize=1 and ParentSnapshotId='ebs-demo'. Below the script, the command 'python3 ebsCreate.py' is executed. The AWS Management Console shows the 'Create Snapshot' action, resulting in a pending snapshot 'snap-0830ade5fff1430b3' of size 1 GiB. The console also displays the details of this snapshot, including its ID, status, volume, and creation time.

```
root@ubuntu: /home/aqc/aqc/lambda/s3

import boto3
import pandas

# Creating the low level functional client
client = boto3.client(
    'ebs',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

# Creating the high level object oriented interface
resource = boto3.resource(
    'ebs',
    aws_access_key_id = 'AKIAQEEOWQHEXQGR5C5C',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

client.start_snapshot(
    VolumeSize=1,
    ParentSnapshotId='ebs-demo'
)

-- INSERT --

root@ubuntu: /home/aqc/aqc/lambda/s3# python3 ebsCreate.py
root@ubuntu: /home/aqc/aqc/lambda/s3#
```

**AWS Management Console - Create Snapshot**

Name	Snapshot ID	Size	Description	Status	Started	Progress	Encryption
	snap-0830ade5fff1430b3	1 GiB		pending	October 27, 2021 at 1:32:57 ...	0%	Not Encrypted

**Snapshot: snap-0830ade5fff1430b3**

Description	
Snapshot ID	snap-0830ade5fff1430b3
Status	pending
Volume	vol-4fffff
Started	October 27, 2021 at 1:32:57 AM UTC-7
Owner	008889205193
Product codes	-
Progress	0%
Capacity	1 GiB
Encryption	Not Encrypted
KMS Key ID	
KMS Key Aliases	
KMS Key ARN	



List snapshot blocks:

```
root@ubuntu: /home/aqc/aqc/lambda/s3
```

```
import boto3
import pandas

# Creating the low level functional client
client = boto3.client(
    'ebs',
    aws_access_key_id = 'AKIAQEEOWQHERBBHGYMW',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'



)

client.list_snapshot_blocks(
    SnapshotId='snap-0920aa5f5c0431927'

)

wq

root@ubuntu:/home/aqc/aqc/lambda/s3# python3 ebsListBlock.py
root@ubuntu:/home/aqc/aqc/lambda/s3#
```

	snap-0920aa5f5c04...	8 GiB	111	 completed	October 27, 2021 at 2:32:38 ...	available (100%)	Not Encrypted
---	----------------------	-------	-----	---	---------------------------------	------------------	---------------

Complete snapshot:

```
root@ubuntu: /home/aqc/aqc/lambda/s3

import boto3
import pandas

# Creating the low level functional client
client = boto3.client(
    'ebs',
    aws_access_key_id = 'AKIAQEEOWQHERBBHGYMW',
    aws_secret_access_key = '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',
    region_name = 'us-west-1'
)

client.complete_snapshot(
    SnapshotId='snap-0920aa5f5c0431927',
    ChangedBlocksCount=0
)

~
~
~
~
~
~
~
~
-- INSERT --
15,2 All

root@ubuntu:/home/aqc/aqc/lambda/s3# python3 ebsCompleteSnapshot.py
root@ubuntu:/home/aqc/aqc/lambda/s3#
```